


```

1 import matplotlib.pyplot as plt
2 import numpy as np

1 x = np.array([0.2,0.6,0.6,1.5,1.8])
2 y = np.array([0.1,0.4,0.9,1.2,1.7])
3 plt.scatter(x,y)

```

⇒ <matplotlib.collections.PathCollection at 0x7f7c812b5210>



```

1 #from sklearn.model_selection import train_test_split
2 #X_train,X_test,y_train,y_test = train_test_split(x,y,test_size=0.6)
3 X_train,y_train = np.array([0.6,1.8]), np.array([0.4,1.7])
4 X_test,y_test = np.array([0.2,0.6,1.5]), np.array([0.1,0.9,1.2])
5 X_train.shape

```

⇒ (2,)

```

1 from sklearn.linear_model import LinearRegression
2 X_train.resize(len(X_train),1)
3 y_train.resize(len(y_train),1)
4 lr = LinearRegression()
5 lr.fit(X_train,y_train)
6 coefs = lr.coef_[0]
7 intercept = lr.intercept_[0]
8 print("y = {:.4f} + {:.4f}x".format(intercept,coefs[0]))
9 X_train.shape

```

⇒ y = -0.2500 + 1.0833x  
(2, 1)


```

1 xx = np.linspace(0,2,2)
2 yy = lr.predict(xx.reshape(len(xx),1))

```

```
3 plt.scatter(x,y)
```

```
4 plt.plot(x_train,y_train,color='red')  
=> [<matplotlib.lines.Line2D at 0x7f7bcb4e2d10>]
```



```
1 from sklearn.metrics import mean_squared_error as mse  
2 mse1 = mse(y_test, lr.predict(X_test.reshape(len(X_test),1)))  
3 print('Error : {}'.format(mse1))
```

```
> Error : 0.09946759259259248
```


```
1
```

```
1 # Regularización Ridge  
2 from sklearn.linear_model import Ridge  
3 ridge = Ridge(alpha = 0.3)  
4 ridge.fit(X_train,y_train)  
5 coefs = ridge.coef_[0]  
6 intercept = ridge.intercept_[0]  
7 print("y = {:.4f} + {:.4f}x".format(intercept,coefs[0]))
```

```
> y = 0.1324 + 0.7647x
```

```
1 xx = np.linspace(0.0,2,2)  
2 yy2 = ridge.predict(xx.reshape(len(xx),1))  
3 plt.scatter(x,y)  
4 plt.plot(xx,yy2,c='g')
```

→ [



```
1 from sklearn.metrics import mean_squared_error as mse
2 mse2 = mse(y_test, ridge.predict(X_test.reshape(len(X_test), 1)))
3 print('Error : {}'.format(mse2))
```

→ Error : 0.04533737024221454


1

```
1 # Regularización Lasso
2 from sklearn.linear_model import Lasso
3 lasso = Lasso(alpha=0.3)
4 lasso.fit(X_train,y_train)
5 coefs = lasso.coef_[0]
6 intercept = lasso.intercept_[0]
7 print("y = {:.4f} + {:.4f}x".format(intercept,coefs))
8 coefs
```

→ y = 0.7500 + 0.2500x  
np.float64(0.25000000000000006)

```
1 xx = np.linspace(0.0,2,2)
2 yy3 = lasso.predict(xx.reshape(len(xx),1))
3 plt.scatter(x,y)
4 plt.plot(xx,yy3,c='g')
```

→ [<matplotlib.lines.Line2D at 0x7f7bcae2e450>]



```
1 from sklearn.metrics import mean_squared_error as mse
2 mse2 = mse(y_test, lasso.predict(X_test.reshape(len(X_test), 1)))
3 print('Error : {}'.format(mse2))
```

→ Error : 0.16520833333333337

1

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 sns.set()
```

```
1 df = pd.read_csv('families.csv')
2 df.tail(2)
```


→

	family	father	mother	midparentHeight	children	childNum	gender	childHeight
932	204	62.5	63.0	65.27	2	1	male	66.5
933	204	62.5	63.0	65.27	2	2	female	57.0

```
1 #import plotly.express as px
2 #fig = px.scatter_3d(df, x='father', y='mother', z='childHeight')
3 #fig.show()
```

```
1 X = df.iloc[:, [1,3]].values
2 y = df.iloc[:, 7].values
```

```
1 plt.plot(X[:,0], y, 'b.')
2 plt.xlabel('Padre')
3 plt.ylabel('Hija(o)')
4 plt.show()
```



```
1 plt.plot(X[:,1], y, 'b.')
2 plt.xlabel('Madre')
3 plt.ylabel('Hija(o)')
4 plt.show()
```



```
1 from sklearn.model_selection import train_test_split
2 from sklearn.linear_model import LinearRegression
3 from sklearn.metrics import r2_score, mean_squared_error
```

```
1 # conjunto de entrenamiento y pruebas
2 X_train,X_test,y_train,y_test=train_test_split(X, y, test_size=0.2,
3                                               random_state=2)
```

```
1 # modelo lineaal y su rendimiento
2 lr = LinearRegression()
3 lr.fit(X_train, y_train)
4 y_pred = lr.predict(X_test)
5 print(' R2 : ',r2_score(y_test, y_pred))
6 print('RMSE : ',np.sqrt(mean_squared_error(y_test, y_pred)))
```



```
R2 : 0.05375000577834388
RMSE : 3.3605503188369497
```

```
1 print(lr.coef_)
2 print(lr.intercept_)
```



```
[0.07286771 0.55935083]
23.14535713417917
```

1

1


1

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt

1 y = [3, 4, 5, 7, 10, 8, 10, 23, 27, 44, 63, 67, 60, 62, 70, 75, 88, 81,
2      87, 95, 108, 135, 151, 169, 179]
3 x = np.arange(len(y))
4 plt.figure(figsize=(6,4))
5 plt.scatter(x, y)
6 plt.show()

```



```

1 # Características polinomiales
2 from sklearn.preprocessing import PolynomialFeatures
3 poly = PolynomialFeatures(degree=2, include_bias=False) # bias en el modelo lineal

```

```

1 # Generación de características nuevas
2 poly_features = poly.fit_transform(x.reshape(-1, 1))
3 poly_features.shape

```

(25, 2)

```

1 # Modelo lineal
2 from sklearn.linear_model import LinearRegression
3 poly_reg_model = LinearRegression()

```

```

1 # Ajuste y predicción
2 poly_reg_model.fit(poly_features, y)
3 y_pred = poly_reg_model.predict(poly_features)

```

```


1 plt.figure(figsize=(6, 4))
2 plt.scatter(x, y)
3 plt.plot(x, y_pred, c="red")
4 plt.title('Regresión polinomial', size=16)

```

```

5 plt.xlabel('x')
6 plt.ylabel('y')
7 plt.show()

```



```

1 # Determinar el grado del polinomio


```

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 x = 6 * np.random.rand(200, 1) - 3
4 y = 0.8*x**2 + 0.9*x + 2 + np.random.randn(200, 1)
5 #ecuación -> y = 0.8x^2 + 0.9x + 2
6 plt.plot(x, y, 'b.')
7 plt.xlabel('x')
8 plt.ylabel('y')
9 plt.show()

```





```
1 from sklearn.model_selection import train_test_split
2 from sklearn.linear_model import LinearRegression
3 from sklearn.preprocessing import PolynomialFeatures
4 from sklearn.metrics import r2_score, mean_squared_error
5 import seaborn as sns
6 sns.set()


1 # conjunto de entrenamiento y pruebas
2 x_train,x_test,y_train,y_test=train_test_split(x, y, test_size=0.2,
3                                               random_state=2)
```

```
1 # modelo lineaal y su rendimiento
2 lr = LinearRegression()
3 lr.fit(x_train, y_train)
4 y_pred = lr.predict(x_test)
5 print(' R2 : ',r2_score(y_test, y_pred))
6 print('RMSE : ',np.sqrt(mean_squared_error(y_test, y_pred)))
```



```
R2 : 0.348786016337134
RMSE : 2.182948912340432
```

```
1 plt.plot(x_train, lr.predict(x_train), color='r')
2 plt.plot(x, y, 'b.')
3 plt.xlabel('x')
4 plt.ylabel('y')
5 plt.show()
```



```

1 # Polinomio de grado 2
2 poly = PolynomialFeatures(degree=2, include_bias=False)
3 x_train_poly = poly.fit_transform(x_train)
4 x_test_poly = poly.transform(x_test)
5 lr = LinearRegression()
6 lr.fit(x_train_poly, y_train)
7 y_pred = lr.predict(x_test_poly)
8 print(' R2 : ',r2_score(y_test, y_pred))
9 print('RMSE : ',np.sqrt(mean_squared_error(y_test, y_pred)))
10 # Visualización
11 """
12 x_new = np.linspace(-3, 3, 200).reshape(200, 1)
13 x_new_poly = poly.transform(x_new)
14 y_new_pred = lr.predict(x_new_poly)
15 plt.plot(x_new, y_new_pred, "r", linewidth=2, label='Curva de ajuste')
16 plt.plot(x_train, y_train, "b.",label='Entrenamiento')
17 plt.plot(x_test, y_test, "g*",label='Pruebas')
18 plt.xlabel("x")
19 plt.ylabel("y")
20 plt.legend()
21 plt.show()
22 """

```



```

R2 : 0.9177735823112784
RMSE : 0.7756885809730303
'\nx_new = np.linspace(-3, 3, 200).reshape(200, 1)\nx_new_poly =
poly.transform(x_new)\ny_new_pred = lr.predict(x_new_poly)\nplt.plot(x_new, y_new_pred,
"r", linewidth=2, label='\Curva de ajuste')\nplt.plot(x_train, y_train,
"b.",label='\Entrenamiento')\nplt.plot(x_test, y_test,
"g*",label='\Pruebas')\nplt.xlabel("x")\nplt.ylabel("y")\nplt.legend()\nplt.show()\n'

```

```

1 # Polinomio de grado 2
2 poly = PolynomialFeatures(degree=2, include_bias=False)

```

```

3 x_train_poly = poly.fit_transform(x_train)
4 x_test_poly = poly.transform(x_test)
5 lr = LinearRegression()
6 lr.fit(x_train_poly, y_train)
7 y_pred = lr.predict(x_test_poly)
8 print(' R2 : ',r2_score(y_test, y_pred))
9 print('RMSE : ',np.sqrt(mean_squared_error(y_test, y_pred)))

```

```

⇒ R2 : 0.9177735823112784
   RMSE : 0.7756885809730303

```

```

1 # Polinomio de grado 3
2 poly = PolynomialFeatures(degree=3, include_bias=False)
3 x_train_poly = poly.fit_transform(x_train)
4 x_test_poly = poly.transform(x_test)
5 lr = LinearRegression()
6 lr.fit(x_train_poly, y_train)
7 y_pred = lr.predict(x_test_poly)
8 print(' R2 : ',r2_score(y_test, y_pred))
9 print('RMSE : ',np.sqrt(mean_squared_error(y_test, y_pred)))

```

```

⇒ R2 : 0.9158615394486871
   RMSE : 0.7846554471212009

```

```

1 # Polinomio de grado 10
2 poly = PolynomialFeatures(degree=10, include_bias=False)
3 x_train_poly = poly.fit_transform(x_train)
4 x_test_poly = poly.transform(x_test)
5 lr = LinearRegression()
6 lr.fit(x_train_poly, y_train)
7 y_pred = lr.predict(x_test_poly)
8 print(' R2 : ',r2_score(y_test, y_pred))
9 print('RMSE : ',np.sqrt(mean_squared_error(y_test, y_pred)))

```

```

⇒ R2 : 0.9008513193120585
   RMSE : 0.8517755762987763

```

1

1

```

1 # PolyReg múltiples características
2 import numpy as np
3 import pandas as pd
4 import matplotlib.pyplot as plt

```

```


1 # Generador de datos
2 np.random.seed(1)
3 x1 = np.absolute(np.random.randn(100, 1) * 10)
4 x2 = np.absolute(np.random.randn(100, 1) * 30)
5 y = 2*x1**2 + 3*x2 + 2 + np.random.randn(100, 1)*20
6 #ecuación -> y = 2*x1^2 + 3*x2 + 2

```

```

1 # Gráficas
2 fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(10, 4))
3 axes[0].scatter(x1, y)
4 axes[1].scatter(x2, y)
5 axes[0].set_title("y = f(x1)")
6 axes[1].set_title("y = f(x2)")
7 plt.show()


```



```

1 # DataFrame de los datos
2 df = pd.DataFrame({'x1':x1.reshape(100,), 'x2':x2.reshape(100,),
3                    'y':y.reshape(100,)}, index=range(0,100))
4 df.tail(2)

```



	x1	x2	y
98	6.200008	24.328550	132.757355
99	6.980320	31.333263	205.167741

```

1 # Características polinomiales y conjuntos de entrenamiento/prueba
2 from sklearn.preprocessing import PolynomialFeatures
3 from sklearn.model_selection import train_test_split
4 X, y = df[['x1', 'x2']], df['y']
5 poly = PolynomialFeatures(degree=2, include_bias=False)
6 X_poly_features = poly.fit_transform(X)
7 X_poly_features.shape

```



(100, 5)

```

1 # Entrenamiento y pruebas
2 X_train,X_test,y_train,y_test=train_test_split(X_poly_features, y,
3                                                  test_size=0.3, random_state=42)

```

```

1 # Modelo lineal
2 from sklearn.linear_model import LinearRegression

```

```

3 poly_reg_model = LinearRegression()

1 # Ajuste, predicción y evaluación
2 poly_reg_model.fit(X_train, y_train)
3 y_pred = poly_reg_model.predict(X_test)
4 from sklearn.metrics import mean_squared_error
5 poly_reg_rmse = np.sqrt(mean_squared_error(y_test, y_pred))
6 poly_reg_rmse

```

```

➡ np.float64(20.937707839078698)

```

```

1 """ Para dos variables independientes se convierte en:
2 indep + a·x1 + b·x2 + c·x1^2+ d·x1·x2 + e·x2^2
3 #ecuación -> y = 2·x1^2 + 3x2 + 2
4 con: indep = intercept
5 a = coefs[0]
6 ...
7 e = coefs[4]
8 """
9 print(poly_reg_model.intercept_, poly_reg_model.coef_)

```

```

➡ 14.123436038978923 [0.61945509 1.9140045 1.89905813 0.0207338 0.01300394]

```

```

1

```

```

1 # Características polinomiales y conjuntos de entrenamiento/prueba
2 from sklearn.preprocessing import PolynomialFeatures
3 from sklearn.model_selection import train_test_split
4 X, y = df[['x1', 'x2']], df['y']

1 # Entrenamiento y pruebas
2 X_train,X_test,y_train,y_test=train_test_split(X, y,
3 test_size=0.3, random_state=42)

```

```

1 # Vs modelo lineal puro
2 lin_reg_model = LinearRegression()
3 lin_reg_model.fit(X_train, y_train)
4 lin_reg_y_pred = lin_reg_model.predict(X_test)
5 lin_reg_rmse = np.sqrt(mean_squared_error(y_test, lin_reg_y_pred))
6 lin_reg_rmse

```

```

➡ np.float64(62.302487453878506)

```

```

1

```

```

1

```

```

1

```

```

1

```

```

1 import pandas as pd
2 import numpy as np
3 # Conjunto Iris
4 # https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data
5 df = pd.read_csv('https://bit.ly/38XWXS4', header=None)
6 df.tail(2)

```



	0	1	2	3	4
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

```

1 # sepal length & petal length de las primeras 100 entradas
2 # Iris-setosa & Iris-versicolor
3 X = df.iloc[:100, [0,2]].values
4 # Cambiando etiquetas de texto a números
5 y = df.iloc[:100, 4].values
6 y = np.array(np.where(y=='Iris-setosa',-1,1))
7 y

```



```


array([-1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1,
       -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1,
       -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1, -1,
        1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,
        1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,
        1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1])

```

```

1 import matplotlib.pyplot as plt
2 plt.scatter(X[:50,0],X[:50,1],color='red',marker='o',label='Setosa')
3 plt.scatter(X[50:,0],X[50:,1],color='blue',marker='^',label='Versicolor')
4 plt.xlabel('sepal-length [cm]')
5 plt.ylabel('petal-length [cm]')
6 plt.legend(loc='upper left')
7 plt.show()

```




```
1 # Perceptrón
2 from sklearn.linear_model import Perceptron
3 ppn = Perceptron(max_iter=40, eta0=0.1, random_state=1)
4 ppn.fit(X, y)
5 print(ppn.intercept_, ppn.coef_)
```



```
[-0.1] [[-0.28  0.7 ]]
```

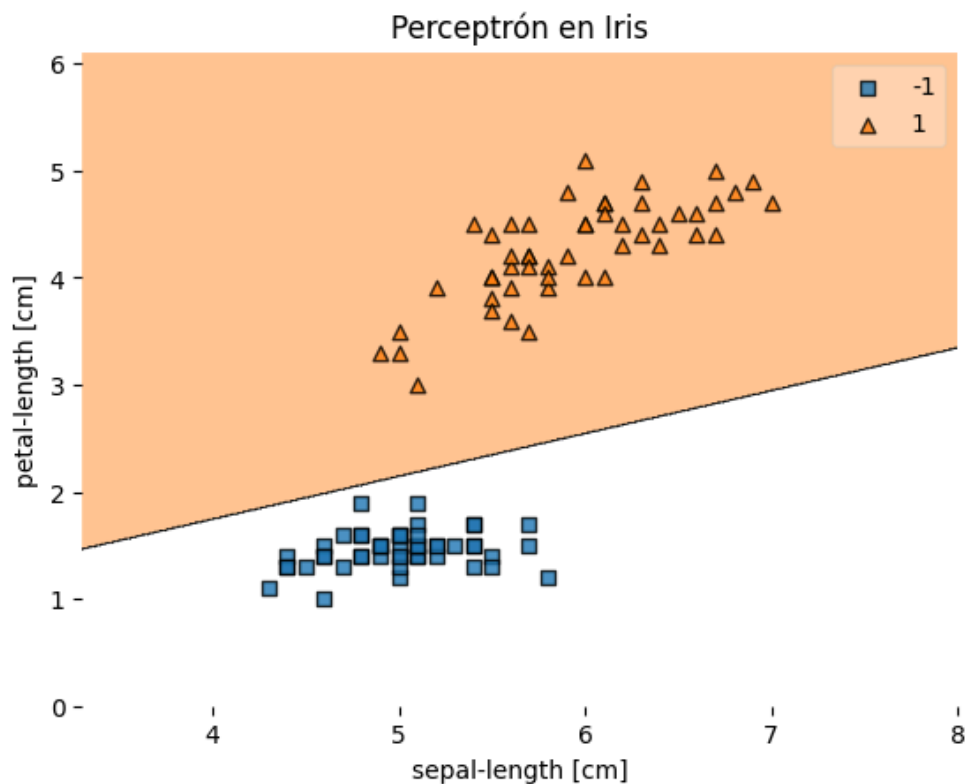
```
1 #      w0 + w1 * x1 + w2 * x2 = 0
2 # intercept_ + coef_[0][0]*x1 + coef_[0][1]*x2
3 # => x2 = -(w0 + w1*x1) / w2
4 # Borde de decision
5 """
6 x1 = np.linspace(4, 7.5, 2)
7 x2 = -(ppn.intercept_ + ppn.coef_[0][0]*x1) / ppn.coef_[0][1]
8 plt.plot(x1, x2, 'g', label = "Borde de Decision")
9 # Clase -1 : setosa
10 plt.scatter(X[y==1][:,0],X[y==1][:,1],color='red',marker='o',label='Setosa')
11 # Clase 1 : versicolor
12 plt.scatter(X[y==0][:,0],X[y==0][:,1],color='blue',marker='^',label='Versicolor')
13 plt.legend()
14 plt.grid(color='lightgray', linestyle='--')
15 """
16 x1 = np.linspace(4, 7.5, 2)
17 x2 = -(ppn.intercept_ + ppn.coef_[0][0]*x1) / ppn.coef_[0][1]
18 plt.plot(x1, x2, 'g', label = "Borde de Decision")
19 # Clase -1 : setosa
20 plt.scatter(X[y==1][:,0],X[y==1][:,1],color='red',marker='o',label='Setosa')
21 # Clase 1 : versicolor
22 plt.scatter(X[y==0][:,0],X[y==0][:,1],color='blue',marker='^',label='Versicolor')
23 plt.legend()
24 plt.grid(color='lightgray', linestyle='--')
```



```
1 #!pip install mlxtend
```

```
1 from mlxtend.plotting import plot_decision_regions
2 # Regiones de decisión
3 plot_decision_regions(X, y, clf=ppn)
4 plt.title('Perceptrón en Iris')
5 plt.xlabel('sepal-length [cm]')
6 plt.ylabel('petal-length [cm]')
7 plt.show()
```





1

```
1 # Con sklearn
2 from sklearn import datasets
3 import numpy as np
4 iris = datasets.load_iris()
```

```
1 X = iris.data[:, [2,3]]
2 y = iris.target[:]
3 print('Etiquetas : ',np.unique(y))
```



Etiquetas : [0 1 2]

```
1 from sklearn.model_selection import train_test_split
2 X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,
3                                               random_state=1,stratify=y)
4 print(X_train.shape, y_train.shape)
5 print(X_test.shape, y_test.shape)
```



(105, 2) (105,)  
(45, 2) (45,)

```
1 print('Total de etiquetas en y      :', np.bincount(y))
2 print('Total de etiquetas en y_train :', np.bincount(y_train))
3 print('Total de etiquetas en y_test  :', np.bincount(y_test))
```



Total de etiquetas en y : [50 50 50]  
Total de etiquetas en y\_train : [35 35 35]  
Total de etiquetas en y\_test : [15 15 15]

```

1 # Estandarización
2 from sklearn.preprocessing import StandardScaler
3 sc = StandardScaler()
4 #sc.fit(X_train)
5 X_train_std = sc.fit_transform(X_train)
6 X_test_std = sc.transform(X_test)

1 # Perceptrón
2 from sklearn.linear_model import Perceptron
3 ppn = Perceptron(max_iter=40, eta0=0.1, random_state=1)
4 # Ajuste y evaluación
5 ppn.fit(X_train_std, y_train)
6 print('Exactitud : ',ppn.score(X_test_std,y_test))

```

⇒ Exactitud : 0.9777777777777777

```

1 from sklearn.metrics import accuracy_score
2 y_pred = ppn.predict(X_test_std)
3 print('Exactitud : ',accuracy_score(y_test,y_pred))

```

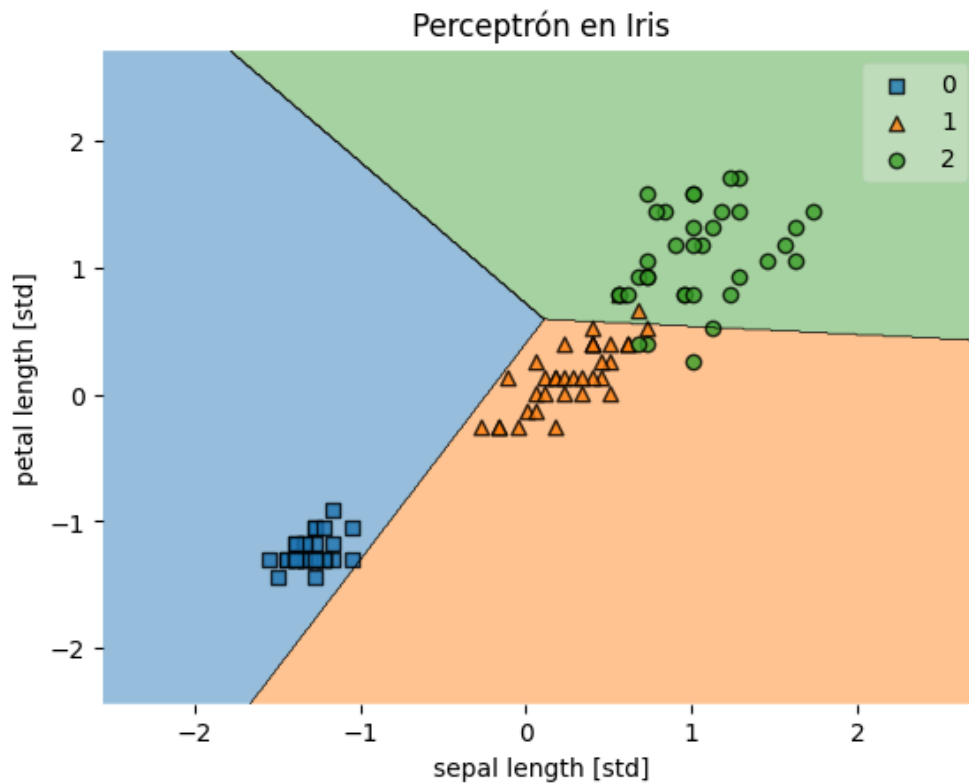
⇒ Exactitud : 0.9777777777777777

```

1 from mlxtend.plotting import plot_decision_regions
2 import matplotlib.pyplot as plt
3 # Regiones de decisión (entrenamiento)
4 plot_decision_regions(X_train_std, y_train, clf=ppn)
5 plt.xlabel('sepal length [std]')
6 plt.ylabel('petal length [std]')
7 plt.title('Perceptrón en Iris')
8 plt.show()

```


⇒



```

1 from mlxtend.plotting import plot_decision_regions
2 import matplotlib.pyplot as plt
3 # Regiones de decisión (pruebas)
4 plot_decision_regions(X_test_std, y_test, clf=ppn)
5 plt.xlabel('sepal length [std]')
6 plt.ylabel('petal length [std]')
7 plt.title('Perceptrón en Iris test')
8 plt.show()

```



1

1 # Regresión Logística

```


1 import numpy as np
2 import matplotlib.pyplot as plt
3 # Sigmoide
4 def sigmoide(z):
5     return 1/(1+np.exp(-z))

```

```

1 z = np.arange(-7,7,0.1)
2 phi_z = sigmoide(z)
3 plt.plot(z, phi_z)
4 plt.axvline(0.0, color='k')
5 plt.ylim(-0.1,1.1)
6 plt.xlabel('z')
7 plt.ylabel('$\phi(z)$')
8 plt.yticks([0.0,0.5,1])
9 ax = plt.gca() # gca = get current axis
10 ax.yaxis.grid(True)
11 plt.show()

```



```
1 from sklearn.linear_model import LogisticRegression
2 lr = LogisticRegression(C=100, random_state=1) # C => inverso de la regularización
3                                             # valor pequeño => regularización más fuerte
4 lr.fit(X_train_std, y_train)
```



```
▼ LogisticRegression ⓘ ?
LogisticRegression(C=100, random_state=1)
```

```
1 y_pred = lr.predict(X_test_std)
2 print('Exactitud : ',lr.score(X_test_std,y_test))
```




```
Exactitud : 0.9777777777777777
```

```
1 # Regiones de decisión (entrenamiento)
2 plot_decision_regions(X_train_std, y_train, clf=lr)
3 plt.xlabel('sepal length [std]')
4 plt.ylabel('petal length [std]')
5 plt.title('Regresión Logística en Iris')
6 plt.show()
```




```
1 # Regiones de decisión (pruebas)
2 plot_decision_regions(X_test_std, y_test, clf=lr)
3 plt.xlabel('sepal length [std]')
4 plt.ylabel('petal length [std]')
5 plt.title('Regresión Logística en Iris')
6 plt.show()
```




1

```
1 # Máquina de soporte vectorial
```

```
1 from mlxtend.plotting import plot_decision_regions
2 import matplotlib.pyplot as plt
3 from sklearn.svm import SVC
4 # Modelo y ajuste
5 svm = SVC(C=0.5, kernel='linear')
6 svm.fit(X_train_std, y_train)
7 # Regiones de decisión (entrenamiento)
8 plot_decision_regions(X_train_std, y_train, clf=svm, legend=2)
9 plt.xlabel('sepal length [std]')
10 plt.ylabel('petal length [std]')
11 plt.title('SVM en Iris')
12 plt.show()
```



```
1 # Regiones de decisión (pruebas)
2 plot_decision_regions(X_test_std, y_test, clf=svm, legend=2)
3 plt.xlabel('sepal length [std]')
4 plt.ylabel('petal length [std]')
5 plt.title('SVM en Iris (pruebas)')
6 plt.show()
```



```

1 # KNN
2 # bibliotecas
3 import numpy as np
4 import pandas as pd
5 import matplotlib.pyplot as plt

```

```

1 data = pd.read_csv('tallas.csv')
2 data.T

```

```

⇒
      0    1    2    3    4    5    6    7    8    9   10   11   12   13   14   15   16   17   18
altura 170 168 163 168 158 160 168 165 160 158 169 158 170 165 161 170 163 160 165 1
peso    64  62  60  63  63  60  66  61  59  59  67  58  63  65  60  68  61  64  62
talla   L   L   M   L   M   M   L   L   M   M   L   M   L   L   M   L   M   L   L

```

```

1 talla_map = {'L':1, 'M':0}
2 data['color'] = data.talla.map(talla_map)
3 data.T

```

```

⇒
      0    1    2    3    4    5    6    7    8    9   10   11   12   13   14   15   16   17   18
altura 170 168 163 168 158 160 168 165 160 158 169 158 170 165 161 170 163 160 165 1
peso    64  62  60  63  63  60  66  61  59  59  67  58  63  65  60  68  61  64  62
talla   L   L   M   L   M   M   L   L   M   M   L   M   L   L   M   L   M   L   L
color    1   1   0   1   0   0   1   1   0   0   1   0   1   1   0   1   0   1   1


```

```

1 plt.scatter
2 new = pd.DataFrame([ [160,63,None,None] ])
3 new.columns = ['altura','peso','talla','color']
4 plt.scatter(new.altura,new.peso,color='r')
5 new2 = pd.DataFrame([ [164,62,None,None] ])
6 new2.columns = ['altura','peso','talla','color']
7 plt.scatter(new2.altura,new2.peso,color='g')
8 plt.scatter(data.altura, data.peso, c=data.color)
9 plt.show()

```





```

1 # knn propio
2 def get_closest_points(data, point, k=3):
3     X = data.iloc[:,0:2].values
4     p = point.iloc[:,0:2].values
5     talla = data.talla.values
6     dist=[[i,np.linalg.norm(X[i]-p),talla[i]]
7           for i in range(len(X))]
8     dist = pd.DataFrame(dist)
9     dist.columns = ['index','dist','talla']
10    return dist.sort_values(by='dist').head(k)
11
12 def show_closest_points(data, point, cercanos, color='k'):
13     plt.scatter(data.altura,data.peso,c=data.color)
14     plt.scatter(point.altura,point.peso,color=color)
15     for c in cercanos.values:
16         p = data.loc[c[0],:]
17         plt.plot([point.altura[0],p.altura],[point.peso[0],p.peso])
18     plt.show()

```

```

1 c = get_closest_points(data, new, 5)
2 print(c)
3 #show_closest_points(data, new, c, color='r')

```



	index	dist	talla
17	17	1.000000	L
4	4	2.000000	M
5	5	3.000000	M
14	14	3.162278	M
19	19	3.162278	L

```

1 c = get_closest_points(data, new2, 5)
2 print(c)
3 #show_closest_points(data, new2, c, color='g')

```

```

⇒ index    dist talla
   18      18  1.000000    L
    7       7  1.414214    L
   16      16  1.414214    M
    2       2  2.236068    M
   19      19  2.236068    L

```

1

```

1 # Con sklearn
2 from sklearn.neighbors import KNeighborsClassifier
3 X = data.iloc[:,2].values
4 y = data.iloc[:,3].values

```

```

1 knn = KNeighborsClassifier(n_neighbors=3)
2 knn.fit(X, y)

```

```

⇒ KNeighborsClassifier ⓘ ?
KNeighborsClassifier(n_neighbors=3)

```

```

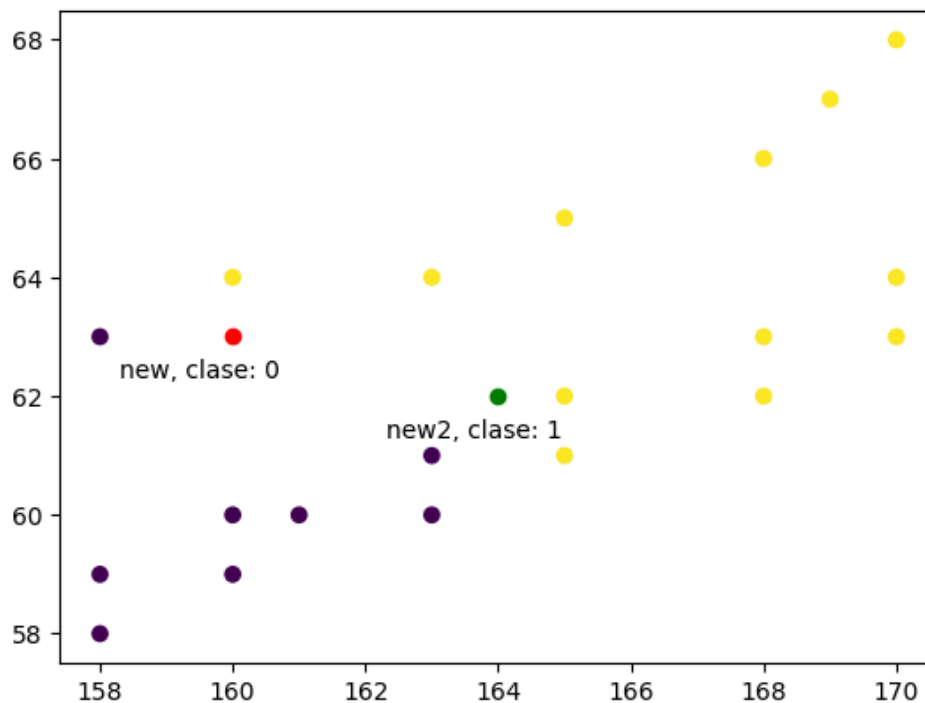
1 new = np.array([160,63]).reshape(1,2)
2 new_pred = knn.predict(new)[0]
3 new2 = np.array([164,62]).reshape(1,2)
4 new2_pred = knn.predict(new2)[0]
5 plt.scatter(X[:,0], X[:,1], c=y)
6 plt.scatter(new[:,0], new[:,1], c='r')
7 plt.text(x=new[:,0]-1.7, y=new[:,1]-0.7, s=f"new, clase: {new_pred}")
8 plt.scatter(new2[:,0], new2[:,1], c='g')
9 plt.text(x=new2[:,0]-1.7, y=new2[:,1]-0.7, s=f"new2, clase: {new2_pred}")

```

```

⇒ Text([162.3], [61.3], 'new2, clase: 1')

```



1

1

1

1

```
1 # Ejemplo 2 ~> regresión
2 import numpy as np
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 import seaborn as sns
```

```
1 # https://archive.ics.uci.edu/dataset/1/abalone
2 url = "https://archive.ics.uci.edu/ml/machine-learning-databases/abalone/abalone.data"
3 abulon = pd.read_csv(url, header=None)
4 abulon.columns = ["Sex", "Length", "Diameter", "Height", "Whole weight",
5                  "Shucked weight", "Viscera weight", "Shell weight", "Rings"]
6 abulon.tail(3)
```




	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.308	9
4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.296	10
4176	M	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.495	12

```
1 # https://laroussecocina.mx/palabra/abulon/
2 # Predecir la edad => No sirve la columna sex
3 abulon = abulon.drop("Sex", axis=1)
```


```
1 sns.pairplot(abulon) # (8x8)
```

 <seaborn.axisgrid.PairGrid at 0x7fbfc87cbd90>




```
1 # número de anillos ~ edad <= Variable objetivo
2 # en este conjunto la mayoría están entre 5 y 15 años
3 abulon.Rings.hist(bins=15)
```

↔ <Axes: >



```
1 #Correlaciones
2 corr = abulon.corr()
3 sns.heatmap(corr, annot=True)
```

↔ <Axes: >



```
1 corr.Rings
```

```
⇒ Length      0.556720
   Diameter    0.574660
   Height      0.557467
   Whole weight 0.540390
   Shucked weight 0.420884
   Viscera weight 0.503819
   Shell weight 0.627574
   Rings       1.000000
   Name: Rings, dtype: float64
```




```
1 X = abulon.drop('Rings',axis=1).values
2 y = abulon.Rings.values
3 X[:,2:],y[2:]
```

```
⇒ (array([[0.455 , 0.365 , 0.095 , 0.514 , 0.2245, 0.101 , 0.15  ],
          [0.35  , 0.265 , 0.09  , 0.2255, 0.0995, 0.0485, 0.07  ]]),
   array([15, 7]))
```

```
1 # Predecir el número de anillo => edad
2 from sklearn.model_selection import train_test_split
3 X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=42)
4 X_train.shape, X_test.shape
```

```
⇒ ((3341, 7), (836, 7))
```

```
1 # Modelo regresor
2 from sklearn.neighbors import KNeighborsRegressor
3 knn = KNeighborsRegressor(n_neighbors=3)
4 knn.fit(X_train, y_train)
```

```
⇒  KNeighborsRegressor  
   KNeighborsRegressor(n_neighbors=3)
```

```
1 from sklearn.metrics import mean_squared_error
2 from math import sqrt
3 y_train_pred = knn.predict(X_train)
4 mse = mean_squared_error(y_train, y_train_pred)
5 sqrt(mse)
```

```
⇒ 1.6643104166366538
```

```
1 y_pred = knn.predict(X_test)
2 mse = mean_squared_error(y_test, y_pred)
3 sqrt(mse)
```

```
⇒ 2.407209400382251
```

```
1
```

```
1
```

```

1 # Ejemplo Bag of Words
2 documents = ['Hello, how are you!',
3             'Win money, win from home.',
4             'Call me now.',
5             'Hello, Call hello you tomorrow?']

```

```

1 # Importar CountVectorizer y crear un objeto
2 from sklearn.feature_extraction.text import CountVectorizer
3 cv = CountVectorizer()
4 cv.fit(documents)
5 names = cv.get_feature_names_out()
6 names

```

```

→ array(['are', 'call', 'from', 'hello', 'home', 'how', 'me', 'money',
        'now', 'tomorrow', 'win', 'you'], dtype=object)

```

```

1 # Transformar y convertir en arreglo
2 docs = cv.transform(documents).toarray()
3 docs

```

```

→ array([[1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1],
        [0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 2, 0],
        [0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0],
        [0, 1, 0, 2, 0, 0, 0, 0, 0, 1, 0, 1]])

```

```

1 # Convertir en DataFrame
2 import pandas as pd
3 freq = pd.DataFrame(data=docs, columns=names)
4 freq

```

```

→
   are  call  from  hello  home  how  me  money  now  tomorrow  win  you
0    1    0    0     1    0    1    0    0    0         0    0    1
1    0    0    1     0    1    0    0    1    0         0    2    0
2    0    1    0     0    0    0    1    0    1         0    0    0
3    0    1    0     2    0    0    0    0    0         1    0    1

```

1

```

1 # Naive Bayes
2 import pandas as pd
3 df = pd.read_csv('spam.csv', names = ['label', 'sms_message'])
4 df.head(3)

```

```

→
   label  sms_message
0    ham  Go until jurong point, crazy.. Available only ...
1    ham                Ok lar... Joking wif u oni...
2   spam  Free entry in 2 a wkly comp to win FA Cup fina...

```

```

1 # Preprocesamiento
2 df.label = df.label.map({'ham':0, 'spam':1})
3 df.head(3)

```

```

⇒

```

	label	sms_message
0	0	Go until jurong point, crazy.. Available only ...
1	0	Ok lar... Joking wif u oni...
2	1	Free entry in 2 a wkly comp to win FA Cup fina...

```

1 from sklearn.model_selection import train_test_split
2 X_train,X_test,y_train,y_test=train_test_split(df.sms_message, df.label,
3                                               random_state=1)
4 print('Conjunto completo: {}'.format(df.shape[0]))
5 print('Conjunto de entrenamiento: {}'.format(X_train.shape[0]))
6 print('Conjunto de pruebas: {}'.format(X_test.shape[0]))

```

```

⇒ Conjunto completo: 5572
Conjunto de entrenamiento: 4179
Conjunto de pruebas: 1393

```

```

1 # Aplicar BoW == Bolsa de palabras
2 from sklearn.feature_extraction.text import CountVectorizer
3 cv = CountVectorizer()
4 # entrenamos el objeto
5 train_data = cv.fit_transform(X_train)
6 test_data = cv.transform(X_test)
7 train_data.shape, test_data.shape

```

```

⇒ ((4179, 7464), (1393, 7464))

```

```

1 # Modelo Bayes multinomial
2 from sklearn.naive_bayes import MultinomialNB
3 nb = MultinomialNB()
4 # Entrenar y predecir
5 nb.fit(train_data, y_train)
6 y_pred = nb.predict(test_data)

```

```

1 # Hay otras metricas
2 from sklearn.metrics import accuracy_score,precision_score,recall_score,f1_score
3 print('Exactitud: ', accuracy_score(y_test,y_pred) )
4 print('Precisión: ', precision_score(y_test,y_pred) )
5 print('Recall: ', recall_score(y_test,y_pred) )
6 print('F1: ', f1_score(y_test,y_pred) )

```

```

⇒ Exactitud: 0.9856424982053122
Precisión: 0.9545454545454546
Recall: 0.9333333333333333
F1: 0.9438202247191011

```



```

1 # XGB Clasificador => Predecir diabetes
2 import numpy as np
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 import seaborn as sns

```

```
1 #!pip install xgboost
```

```

1 from xgboost import XGBClassifier
2 from sklearn.model_selection import train_test_split
3 from sklearn.metrics import accuracy_score

```

```

1 # Datos
2 df = pd.read_csv('diabetes.csv')
3 df.tail(2)

```

```

⇒
      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  DiabetesPedigreeFunction
766             1     126             60             0           0  30.1                0.357
767             1      93             70             31           0  30.4                0.368

```

```

1 # Separar X, y
2 X = df.iloc[:, :-1].values
3 y = df.iloc[:, -1].values

```

```

1 # Entrenamiento y prueba
2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=5)
3 X_train.shape, X_test.shape

```

```
⇒ ((614, 8), (154, 8))
```

```

1 # Modelo
2 xgb = XGBClassifier()
3 xgb.fit(X_train, y_train)

```

```

⇒
XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              feature_weights=None, gamma=None, grow_policy=None,
              importance_type=None, interaction_constraints=None,
              learning_rate=None, max_bin=None, max_cat_threshold=None,
              max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
              max_leaves=None, min_child_weight=None, missing=nan,
              monotone_constraints=None, multi_strategy=None, n_estimators=None,
              n_jobs=None, num_parallel_tree=None, ...)

```

```
1 y_pred = xgb.predict(X_test)
```

```
1 accuracy = accuracy_score(y_test, y_pred)
2 print("Accuracy : ",accuracy)
```

```
➡ Accuracy : 0.7922077922077922
```

```
1
```

```
1 #!pip install xgboost
```

```
1 # Regresor => Predecir calorías
2 import numpy as np
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6 from sklearn.model_selection import train_test_split
7 from xgboost import XGBRegressor
8 from sklearn import metrics
```

```
1 # Datos
2 exercise = pd.read_csv('exercise.csv')
3 exercise.tail(3)
```

```
➡
```

	User_ID	Gender	Age	Height	Weight	Duration	Heart_Rate	Body_Temp
<b>14997</b>	17271188	female	43	159.0	58.0	16.0	90.0	40.1
<b>14998</b>	18643037	male	78	193.0	97.0	2.0	84.0	38.3
<b>14999</b>	11751526	male	63	173.0	79.0	18.0	92.0	40.5

```
1 # Datos
2 calories = pd.read_csv('calories.csv')
3 calories.tail(3)
```

```
➡
```

	User_ID	Calories
<b>14997</b>	17271188	75.0
<b>14998</b>	18643037	11.0
<b>14999</b>	11751526	98.0

```
1 # Combinando los dos DFs
2 calories_data = pd.concat([exercise, calories.Calories], axis=1)
3 calories_data.head(3)
```

```
➡
```

	User_ID	Gender	Age	Height	Weight	Duration	Heart_Rate	Body_Temp	Calories
<b>0</b>	14733363	male	68	190.0	94.0	29.0	105.0	40.8	231.0
<b>1</b>	14861698	female	20	166.0	60.0	14.0	94.0	40.3	66.0
<b>2</b>	11179863	male	69	179.0	79.0	5.0	88.0	38.7	26.0

```
1 # Número de filas y columnas
2 calories_data.shape
```

```
⇒ (15000, 9)
```

```
1 # Información del conjunto
2 calories_data.info()
```

```
⇒ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 15000 entries, 0 to 14999
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  -
0   User_ID     15000 non-null  int64
1   Gender      15000 non-null  object
2   Age         15000 non-null  int64
3   Height      15000 non-null  float64
4   Weight      15000 non-null  float64
5   Duration    15000 non-null  float64
6   Heart_Rate  15000 non-null  float64
7   Body_Temp   15000 non-null  float64
8   Calories    15000 non-null  float64
dtypes: float64(6), int64(2), object(1)
memory usage: 1.0+ MB
```

```
1 # ¿Hay nulos?
2 calories_data.isnull().sum()
```

```
⇒ User_ID      0
   Gender      0
   Age         0
   Height      0
   Weight      0
   Duration    0
   Heart_Rate  0
   Body_Temp   0
   Calories    0
dtype: int64
```

```
1 # Algunas estadísticas de las columnas
2 calories_data.describe()
```

```
⇒
```

	User_ID	Age	Height	Weight	Duration	Heart_Rate	Body_Temp
<b>count</b>	1.500000e+04	15000.000000	15000.000000	15000.000000	15000.000000	15000.000000	15000.000000
<b>mean</b>	1.497736e+07	42.789800	174.465133	74.966867	15.530600	95.518533	40.025453
<b>std</b>	2.872851e+06	16.980264	14.258114	15.035657	8.319203	9.583328	0.779230
<b>min</b>	1.000116e+07	20.000000	123.000000	36.000000	1.000000	67.000000	37.100000
<b>25%</b>	1.247419e+07	28.000000	164.000000	63.000000	8.000000	88.000000	39.600000
<b>50%</b>	1.499728e+07	39.000000	175.000000	74.000000	16.000000	96.000000	40.200000
<b>75%</b>	1.744928e+07	56.000000	185.000000	87.000000	23.000000	103.000000	40.600000
<b>max</b>	1.999965e+07	79.000000	222.000000	132.000000	30.000000	128.000000	41.500000

```
1 # Visualizaciones
2 sns.set()
```

```
1 # Conteo por género
2 #sns.countplot(x=calories_data.Gender)
```

```
1 # Distribución de edades
2 #sns.kdeplot(calories_data.Age, fill=True)
```


```
1 # Histograma de edades
2 #sns.histplot(calories_data.Age, kde=True)
```

```
1 # # Distribución de Peso
2 #sns.histplot(calories_data.Weight, kde=True)
```

```
1 # Correlación entre las columnas (positiva / negativa)
2 corr = calories_data.drop(['Gender'],axis=1).corr()
3 #corr
```

```
1 # Visualización con un heatmap
2 plt.figure(figsize=(10,10))
3 sns.heatmap(corr, cbar=True, annot=True, fmt='.3f',
4             annot_kws={'size':8}, cmap='Blues')
```

[↕] <Axes: >



```
1 # Convertir texto a número
2 calories_data.replace({ 'Gender':{'male':0,'female':1} },inplace=True)
3 calories_data.head()
```

```
→ /tmp/ipykernel_660/941247895.py:2: FutureWarning: Downcasting behavior in `replace` is deprecated
calories_data.replace({'Gender':{'male':0,'female':1}}, inplace=True)
```

	User_ID	Gender	Age	Height	Weight	Duration	Heart_Rate	Body_Temp	Calories
0	14733363	0	68	190.0	94.0	29.0	105.0	40.8	231.0
1	14861698	1	20	166.0	60.0	14.0	94.0	40.3	66.0
2	11179863	0	69	179.0	79.0	5.0	88.0	38.7	26.0
3	16180408	1	34	179.0	71.0	13.0	100.0	40.5	71.0
4	17771927	1	27	154.0	58.0	10.0	81.0	39.8	35.0

```
1 # Separar predictoras y objetivo
2 X = calories_data.drop(columns=['User_ID','Calories'], axis=1)
3 y = calories_data.Calories
```

```
1 # Entrenamiento y pruebas
2 X_train,X_test,y_train,y_test=train_test_split(X, y, test_size=0.2,
3                                               random_state=2)
```

```
1 X.shape, X_train.shape, X_test.shape
```

```
→ ((15000, 7), (12000, 7), (3000, 7))
```

```
1 # Modelo XGBRegressor
2 xgb = XGBRegressor()
3 xgb.fit(X_train, y_train)
```

```
→ XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              feature_weights=None, gamma=None, grow_policy=None,
              importance_type=None, interaction_constraints=None,
              learning_rate=None, max_bin=None, max_cat_threshold=None,
              max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
              max_leaves=None, min_child_weight=None, missing=nan,
              monotone_constraints=None, multi_strategy=None, n_estimators=None,
              n_jobs=None, num_parallel_tree=None, ...)
```

```
1 # Evaluación
2 y_hat = xgb.predict(X_test)
```

```
1 # Error absoluto medio
2 mae = metrics.mean_absolute_error(y_test, y_hat)
3 print("Mean Absolute Error = ", mae)
```

```
→ Mean Absolute Error = 1.4833678883314132
```